

张振，KAFKA MEETUP

Jitney, Kafka at Airbnb



关于Airbnb爱彼迎

成立于2008年8月，爱彼迎总部位于加利福尼亚州旧金山市。爱彼迎是一个值得信赖的社区型市场，在这里人们可以通过网站、手机或平板电脑发布、发掘和预订世界各地的独特房源。



房客总数

超过
150,000,000个



城市

超过65,000个



城堡

超过1,400个



国家

超过191个



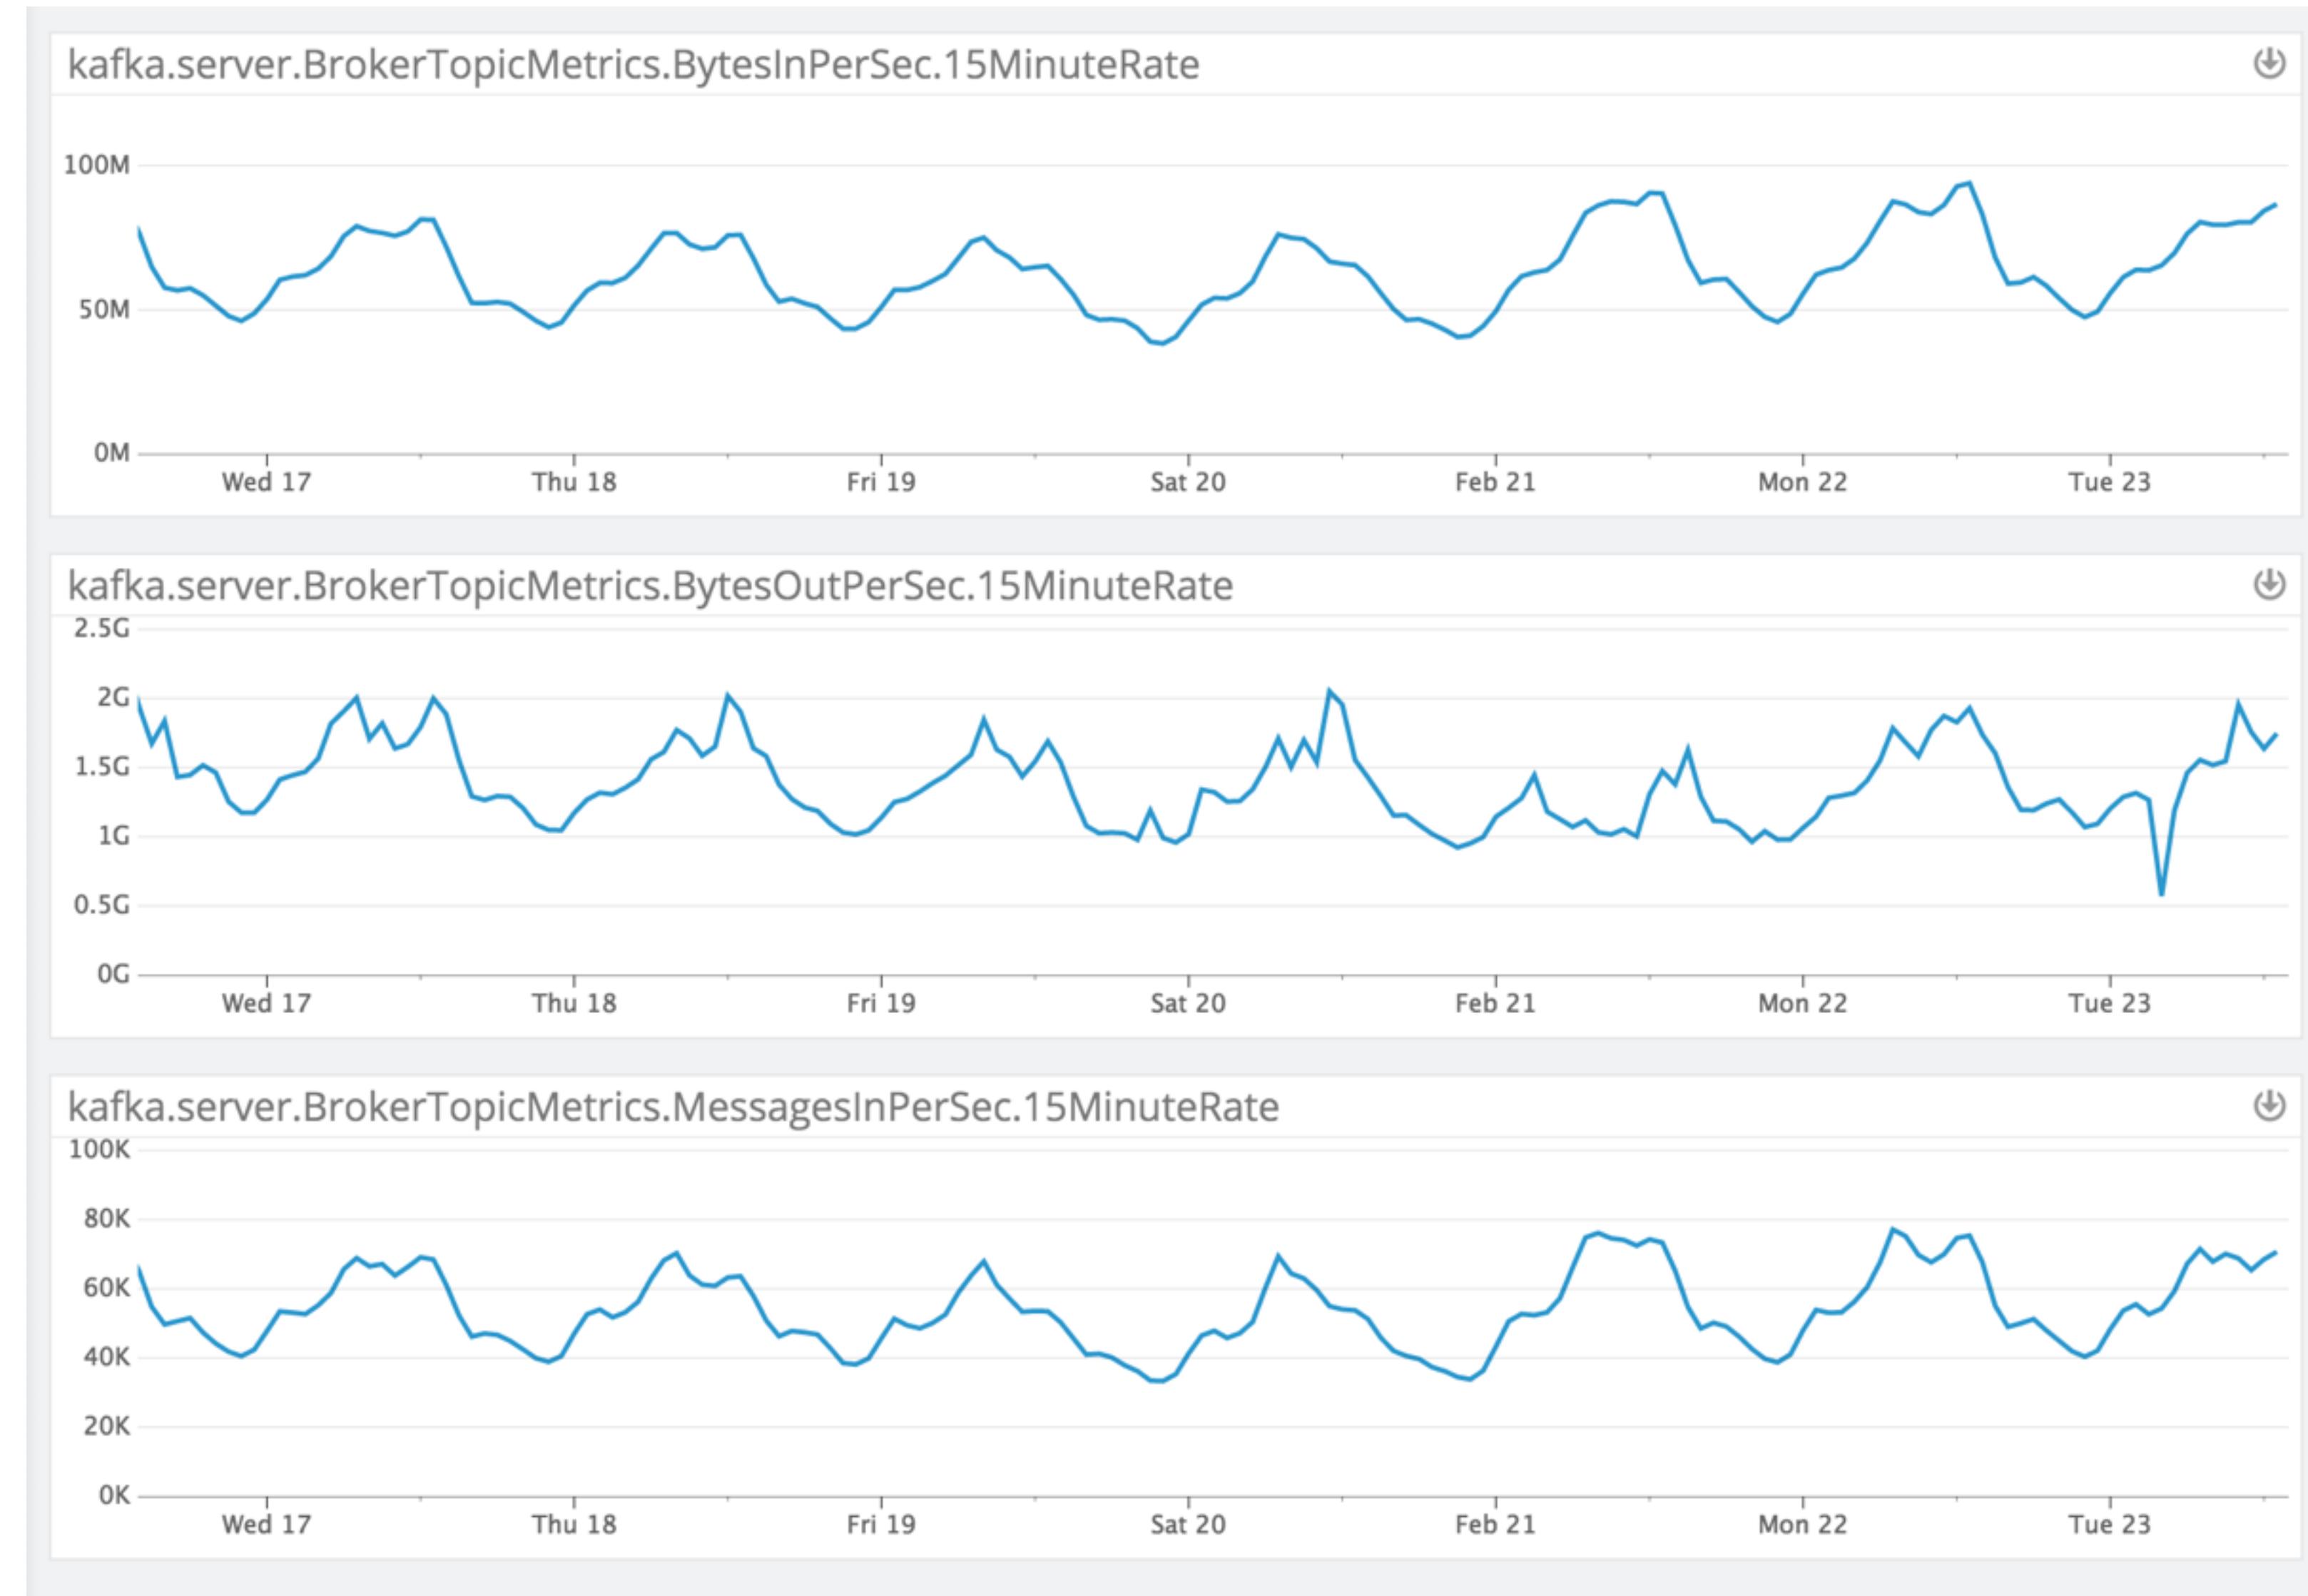
Jitney ?!

a bus carrying passengers for
a low fare



Some Kafka Facts

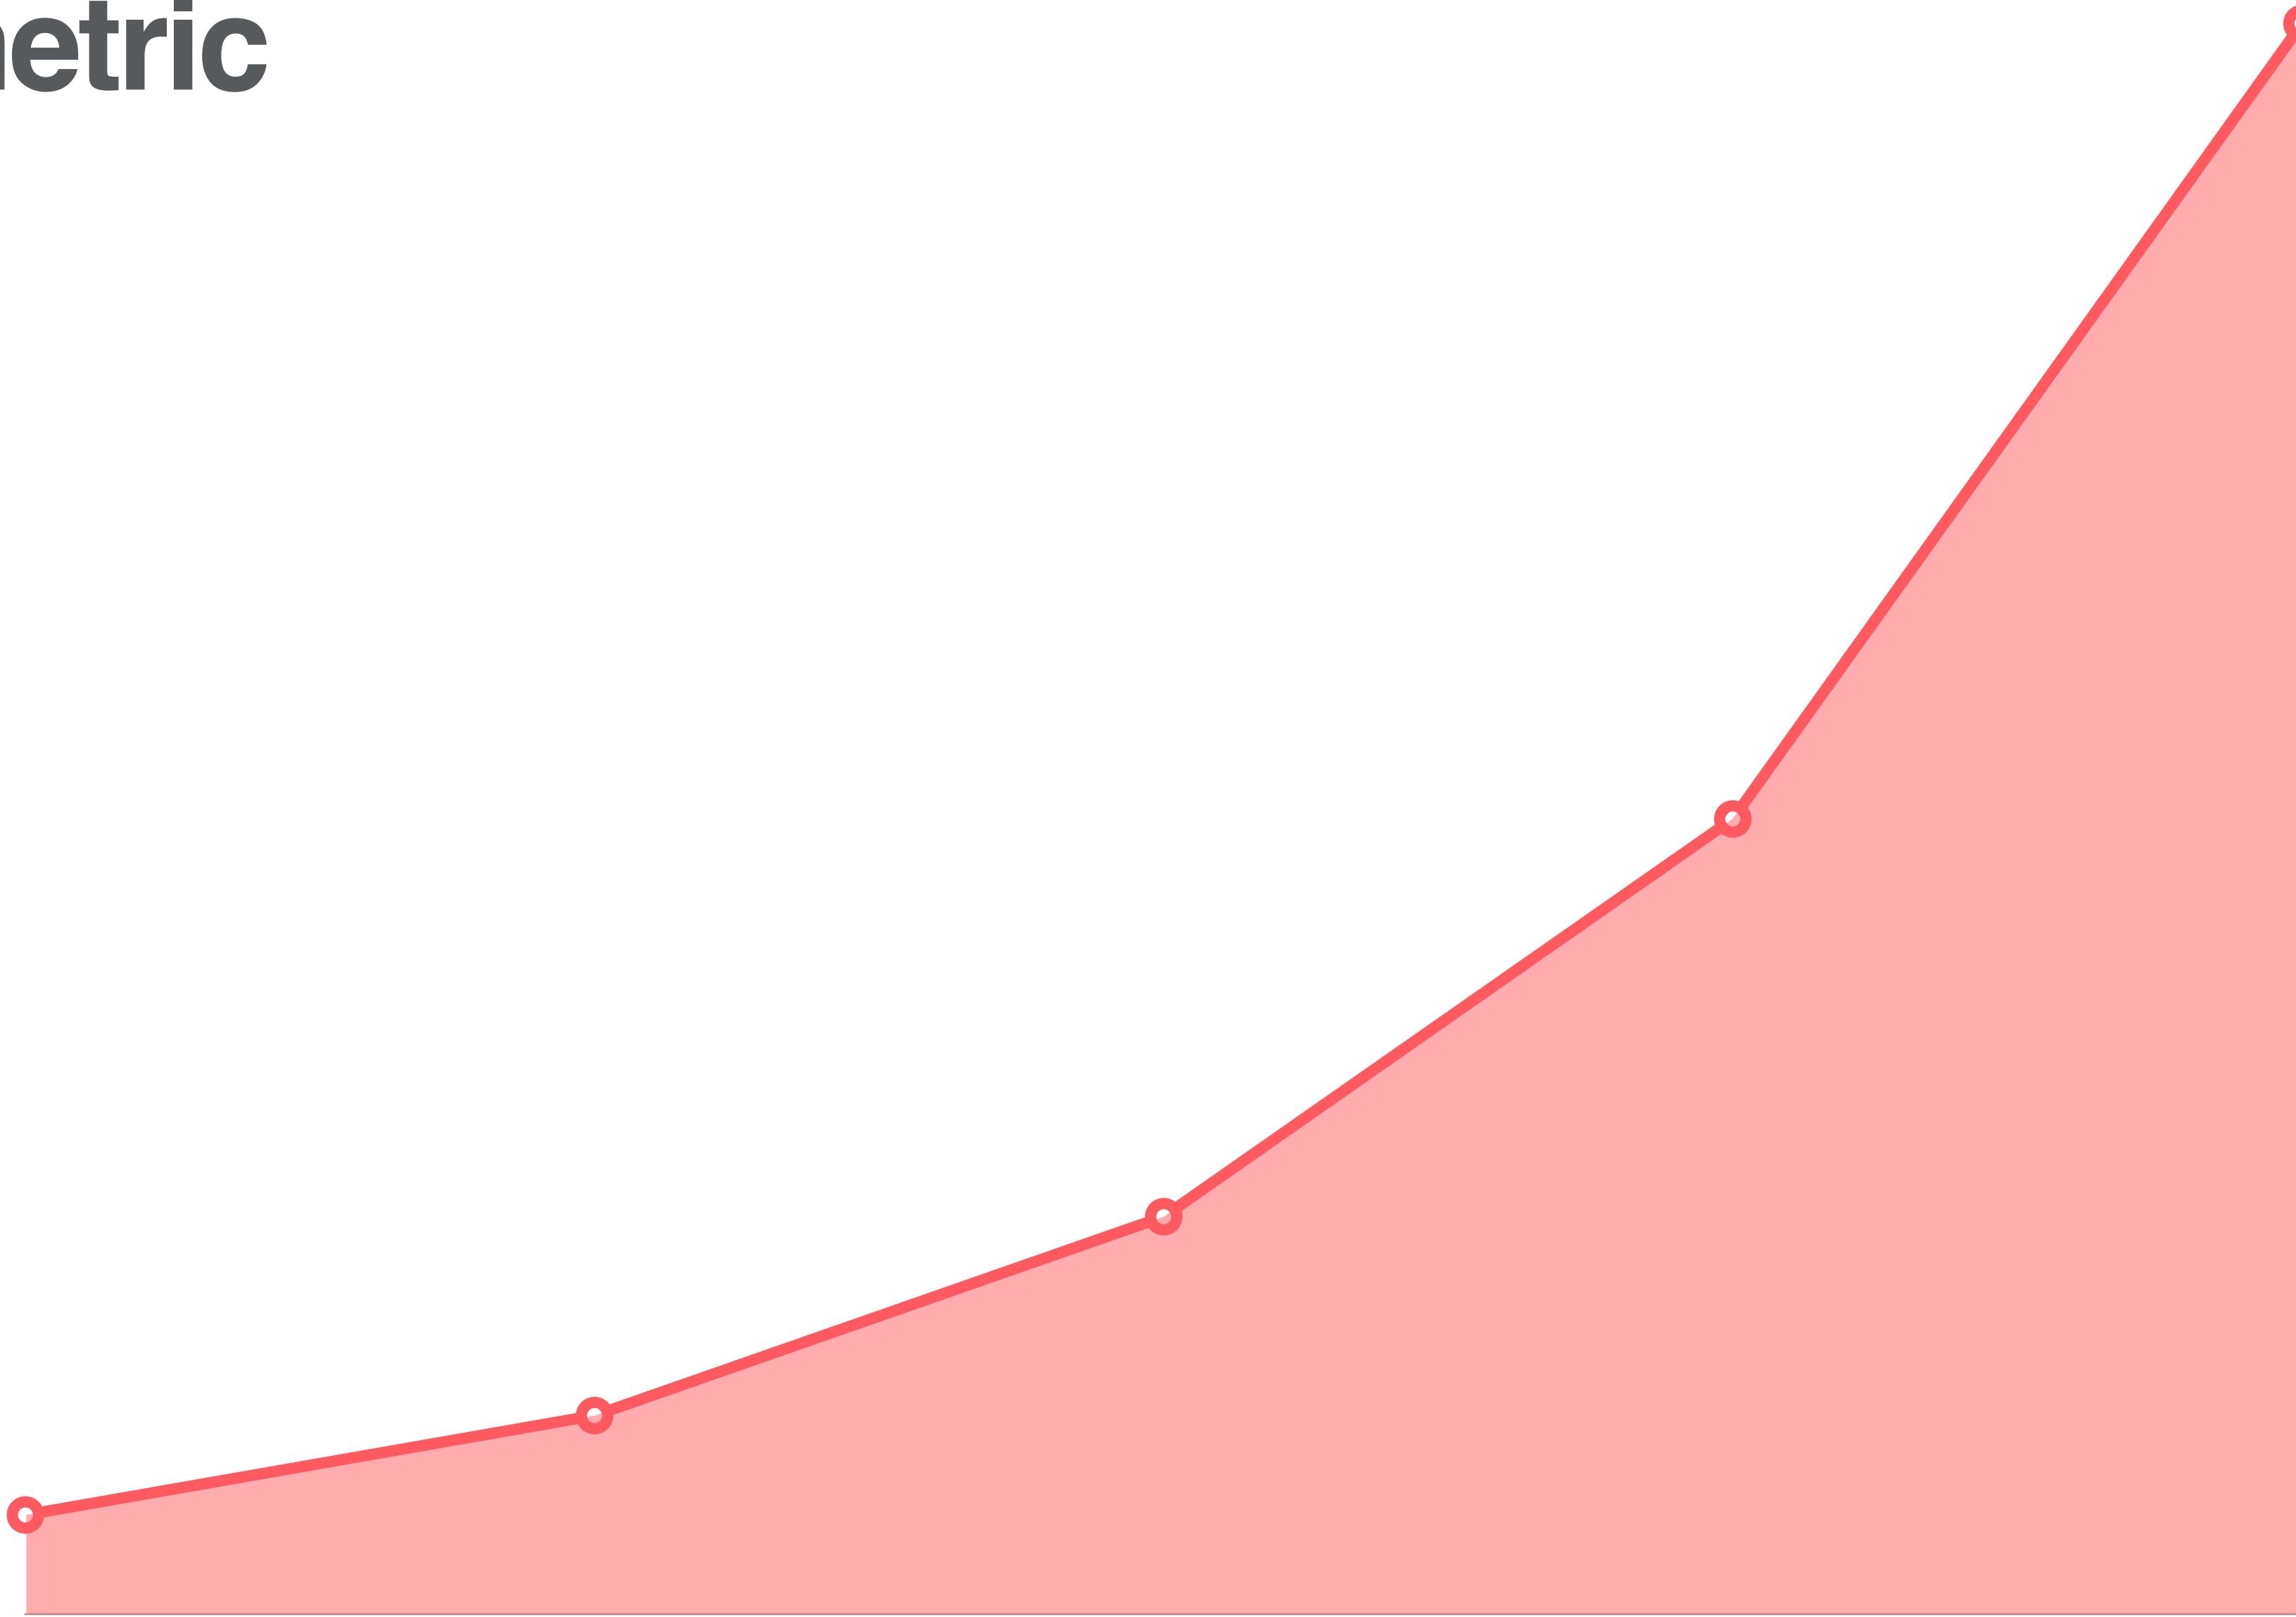
- 1 Production cluster
- v0.8.2
- 90 “small” brokers, d2.2xlarge
- 70 topics
- Replication Factor of 3
- 5 Billions events / day
- IN: 80MB / second
- OUT: 1.5GB / second
- Network bound
- Super stable



A red-tinted photograph of a subway station platform. In the center, a train with red and white stripes is stopped. To the left, a sign reads "JITNEY". The platform has a red and white striped pattern, and there are yellow arrows pointing towards the train.

Why Jitney?

Pick any metric



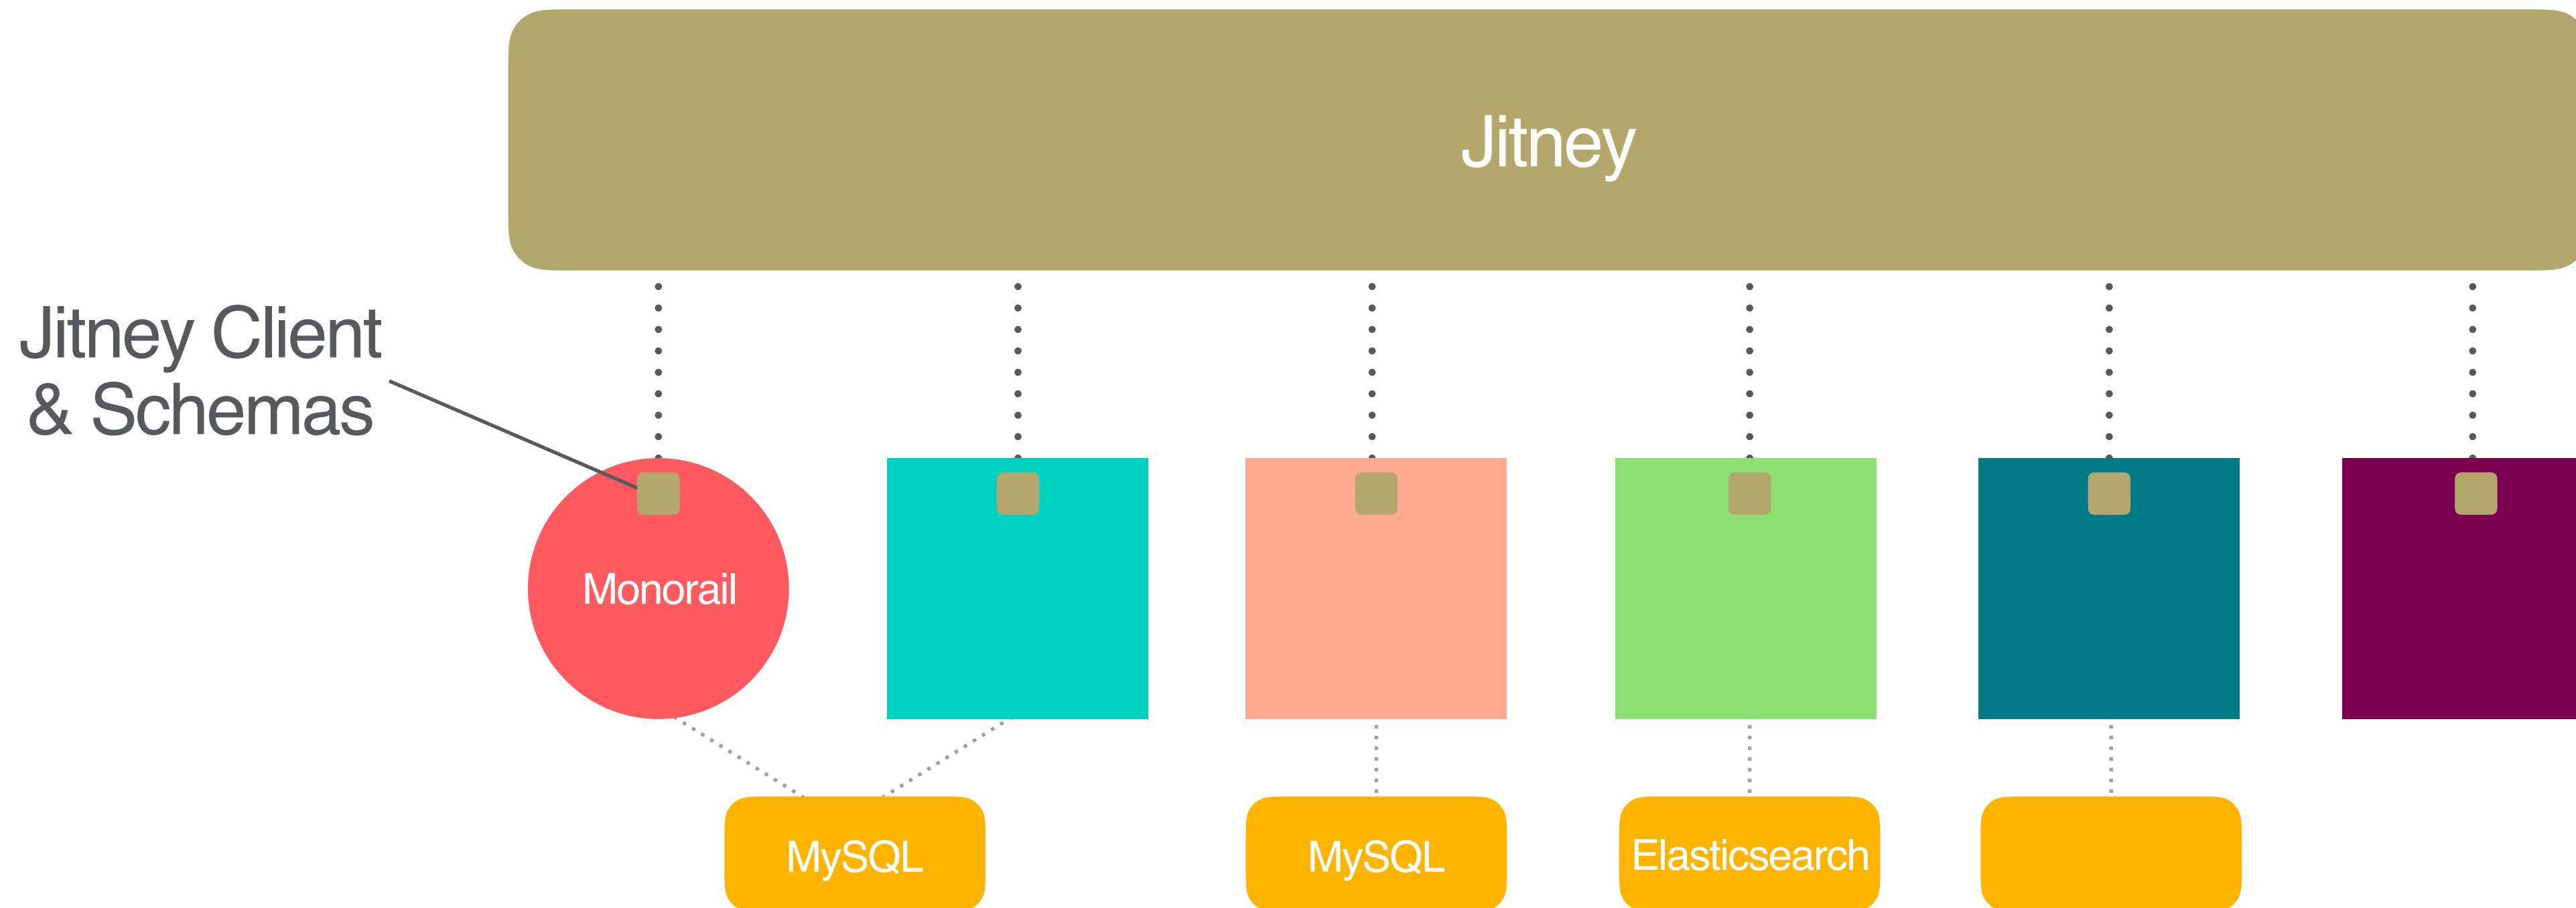
Standardization!





What Use Cases ?

Classic Message Bus



Message Bus

- Decouple Services
- Standard Events
- At-least once delivery
- Standard clients, for Java and Ruby
- Easy to use
- Conventions over configuration

User Activity Logging

- Site and image load times, OOM events
 - Searches, requests, bookings, etc.
 - Experiment assignments
-
- Event data is critical for building data products
 - Data ingestion should be reliable: timely and complete

Challenges

- **JSON events without schemas**
 - **Easy to break events during evolution/code changes**
 - **One topic overall for 800+ event types**
 - **Improper producer configs**
 - **Lack of monitoring**
-
- **Lead to:**
 - **Too many data outages, data loss incidents**
 - **Lack of trust on data systems**

Data Stability

A Year Ago

CEO dashboard and
Magical booking
dashboards were
regularly broken.



Data Stability

A Year Ago

ERF was unstable and experimentation culture was weak

Hi team,

This is partly a PSA to let you know ERF dashboard data hasn't been up to date/accurate for several weeks now. **Do not rely on the ERF dashboard for information about your experiment.**

A vintage white bus with "JITNEY" and "PARK & RIDE" signs on its side, parked on a city street.

Jitney Components

JITNEY BUS
EX 5000
PUBLIC TO REPORT TO CALTRAIN
MARKET AT MISSION AT
PAC-SULL PARK, SEPARATE

Jitney Components

Schema
Repository

Thrift Schema Repository

Why Thrift?

- Easy syntax
- Good performance in Ruby
- Ubiquitous

Advantages of schema repo?

- Great Catalyst for communication, documentation, etc
- it ships jar and gems
- Will developers hate you for this? no



Schema Evolution

- Standard Field in the event schema
- Managed Explicitly
- use Semantic Versioning:

1.0.0 = MODEL . REVISION . ADDITION

MODEL is a change which breaks the rules of backward compatibility.

Example: changing the type of a field.

REVISION is a change which is backward compatible but not forward compatible.

Example: adding a new field to a union type.

ADDITION is a change which is both backward compatible and forward compatible.

Example: adding a new optional field.

Example of Thrift Event

```
# package names must include the major version
namespace java com.airbnb.jitney.event.di_canary.v1
namespace rb airbnb.jitney.event.di_canary.v1

struct CanaryEvent {
    31337: optional string schema = "com.airbnb.jitney.event.di_canary:CanaryEvent:1.0.0"

    1: required i64 canary_created_at,
    2: required string canary_role,
    3: required string canary_seqid,
    4: required i64 canary_seqn,
    5: optional map<string, string> extras,
}
```

Jitney Components

Schema
Repository

Topic Repository

Topic Repository

- **Declare all Jitney topics**
- Aggregate all **characteristics** of a topic:

name

ordering (partitioning function)

whitelist of accepted schemas

- Great for documentation purposes
- DRY

Example of a Topic

```
JitneyTopic.builder()
    .withRouter(SpinaltapHelper::getRouter)
    .setName("spinaltap")
    .addWhitelistedSchemas(
        "com.airbnb.jitney.event.spinaltap.v1",
        "com.airbnb.jitney.event.spinaltap.v2"
    )
    .withKeyProvider(SpinaltapHelper::getPartitionKey)
    .build()
```

Jitney Components

Schema
Repository

Topic Repository

Clients

Jitney Clients

- **Kafka clients** are hard to use correctly
- it's better with 0.9
- **Committing offsets** is tricky, someone will get it wrong
- even with 0.9
- **Configuration** is a mess

Jitney Clients

it provides:

- metrics reporting: github.com/airbnb/kafka-statsd-metrics2
- configuration for default clusters
- built-in support for Schema Repository and Topic Repository

Consumer:

- offset management to implement at-least once delivery
- polymorphic dispatching to event handler



Example of a Java Producer

```
55 JitneyProducer producer = JitneyProducerFactory.create(  
56     new ProducerConfig("my-client-id"),  
57     new JitneyBusConfig("localhost"),  
58     JitneyRepository.getTopic("canary")  
59 );  
60  
61 producer.startAsync();  
62  
63 // publish  
64 final TBase event = new CanaryEvent(...);  
65 producer.publish(event);  
66  
67 producer.stopAsync();
```

Example of a Java Consumer

```
54
55    final Dispatcher dispatcher = ClassBasedDispatcher.builder()
56        .addHandler(CanaryEvent.class, new Handler<CanaryEvent>(){
57            public void handle(Message<CanaryEvent> message) {
58                System.out.println("message received!");
59            }
60        })
61        .build();
62
63    consumer = JitneyConsumerFactory.create(
64        ...,
65        ...,
66        Arrays.asList(Subscription.of(Route.of("canary"), dispatcher)));
67
68
69    consumer.startAsync();
70    // worker threads are created and process incoming messages
```

Jitney Components

Schema
Repository

Topic Repository

Clients

HTTP Proxy

Jitney Components

**Schema
Repository**

Topic Repository

Clients

HTTP Proxy

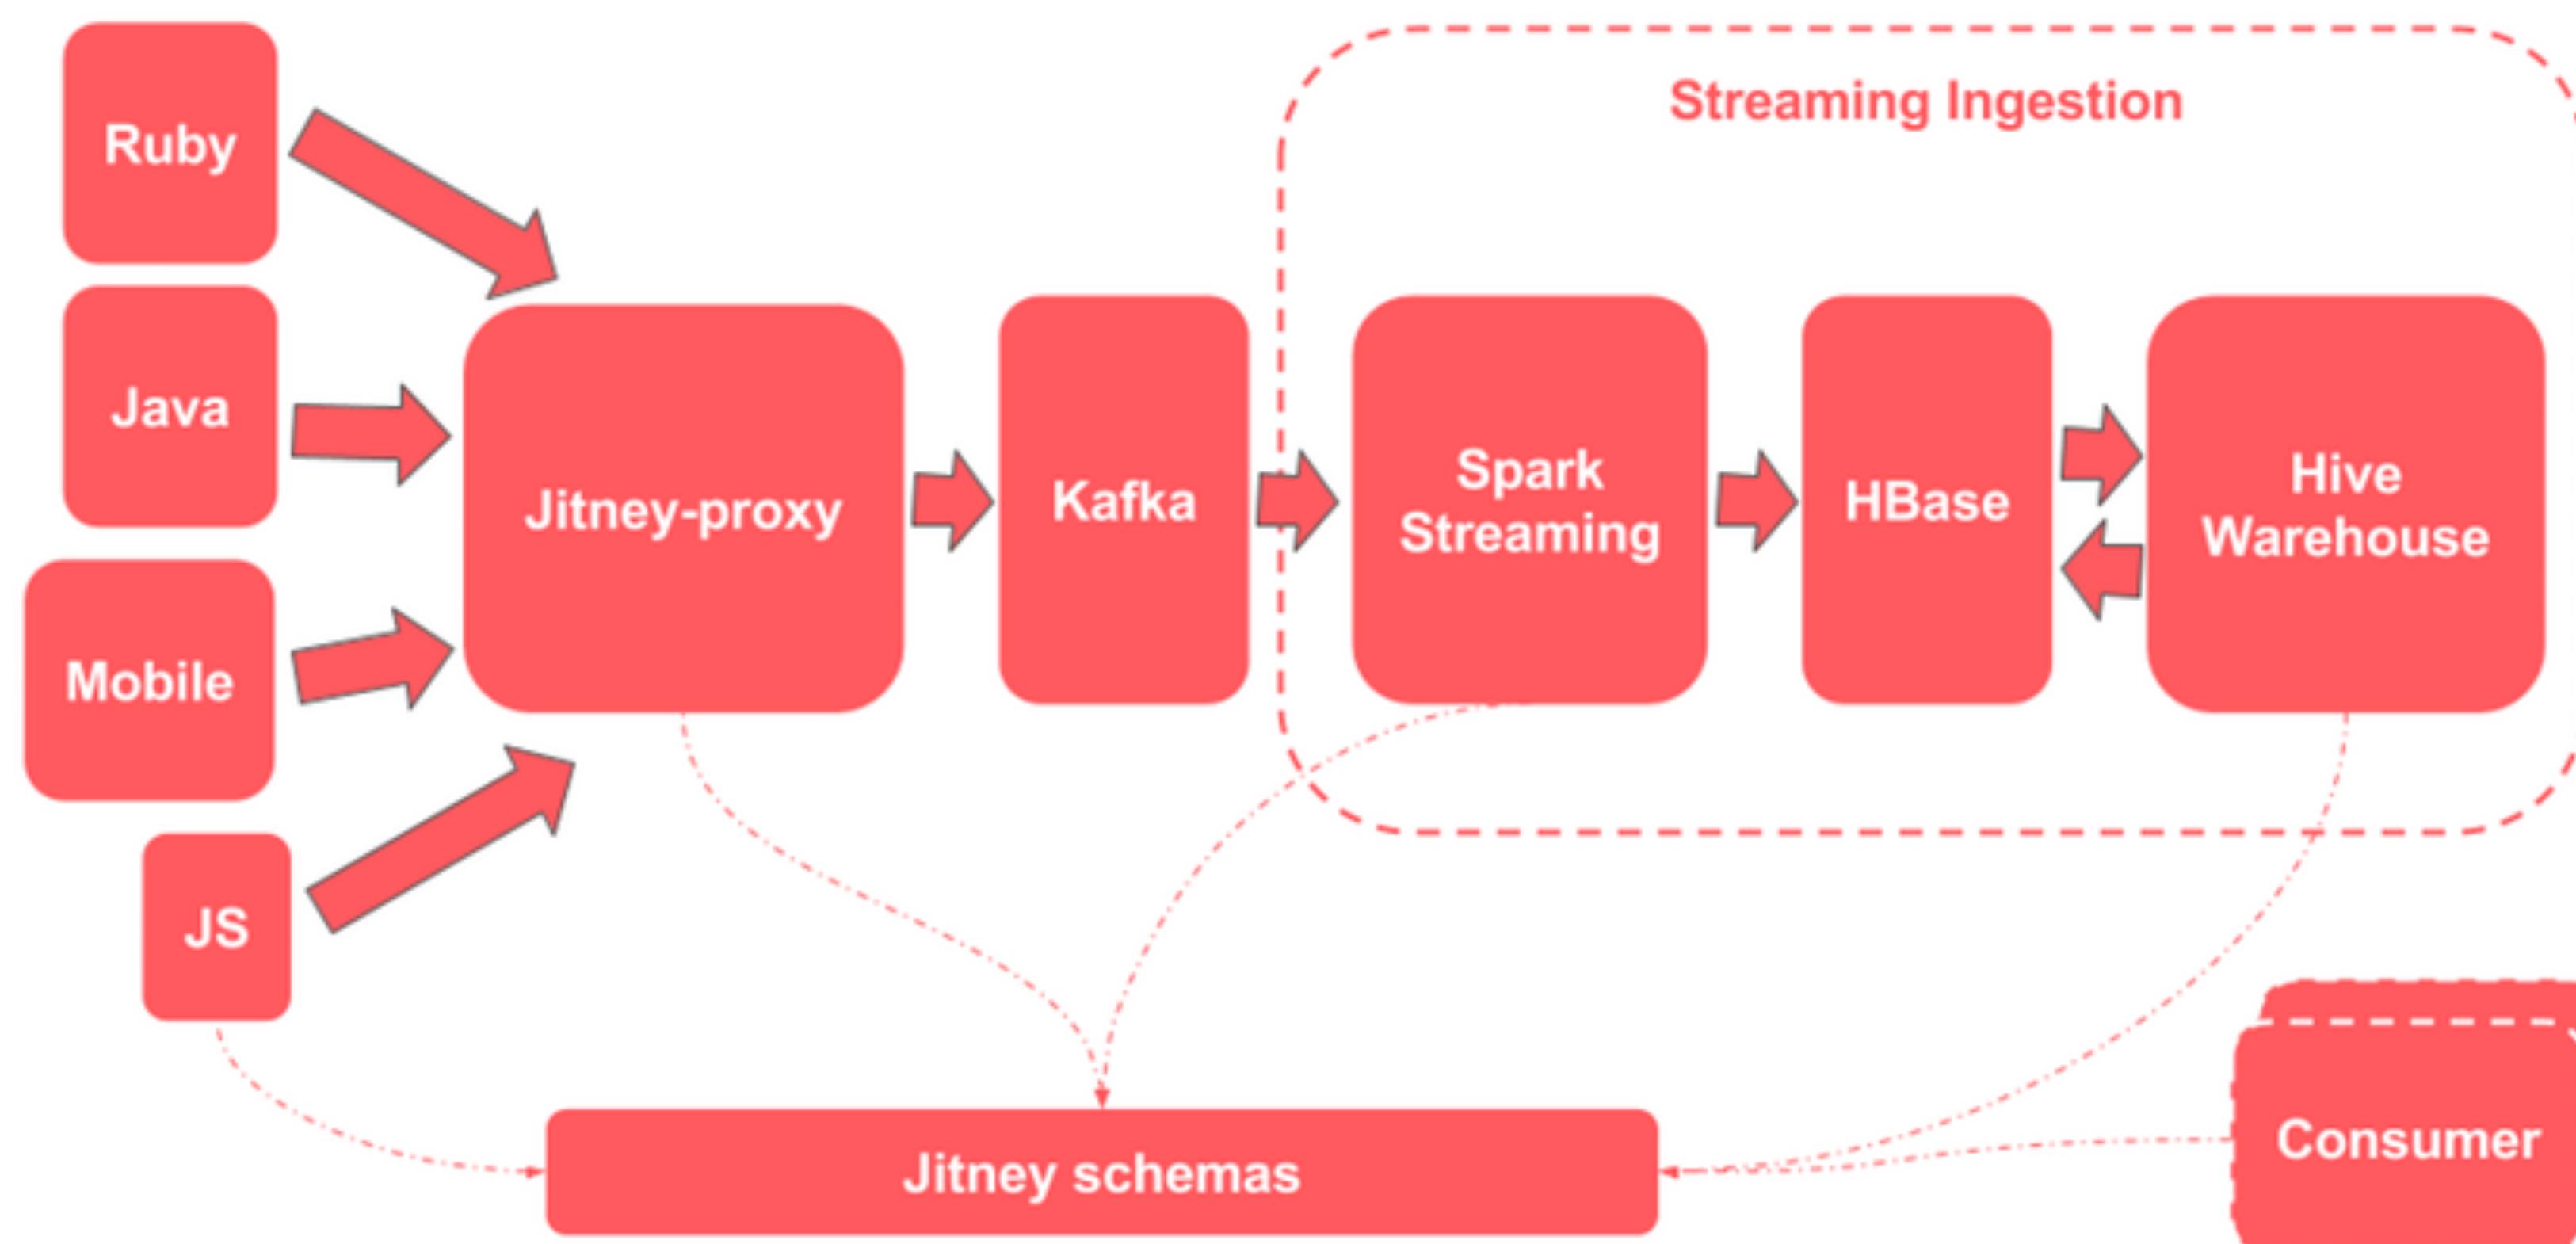
**Warehouse
Integration**

Data Ingestion Pipeline

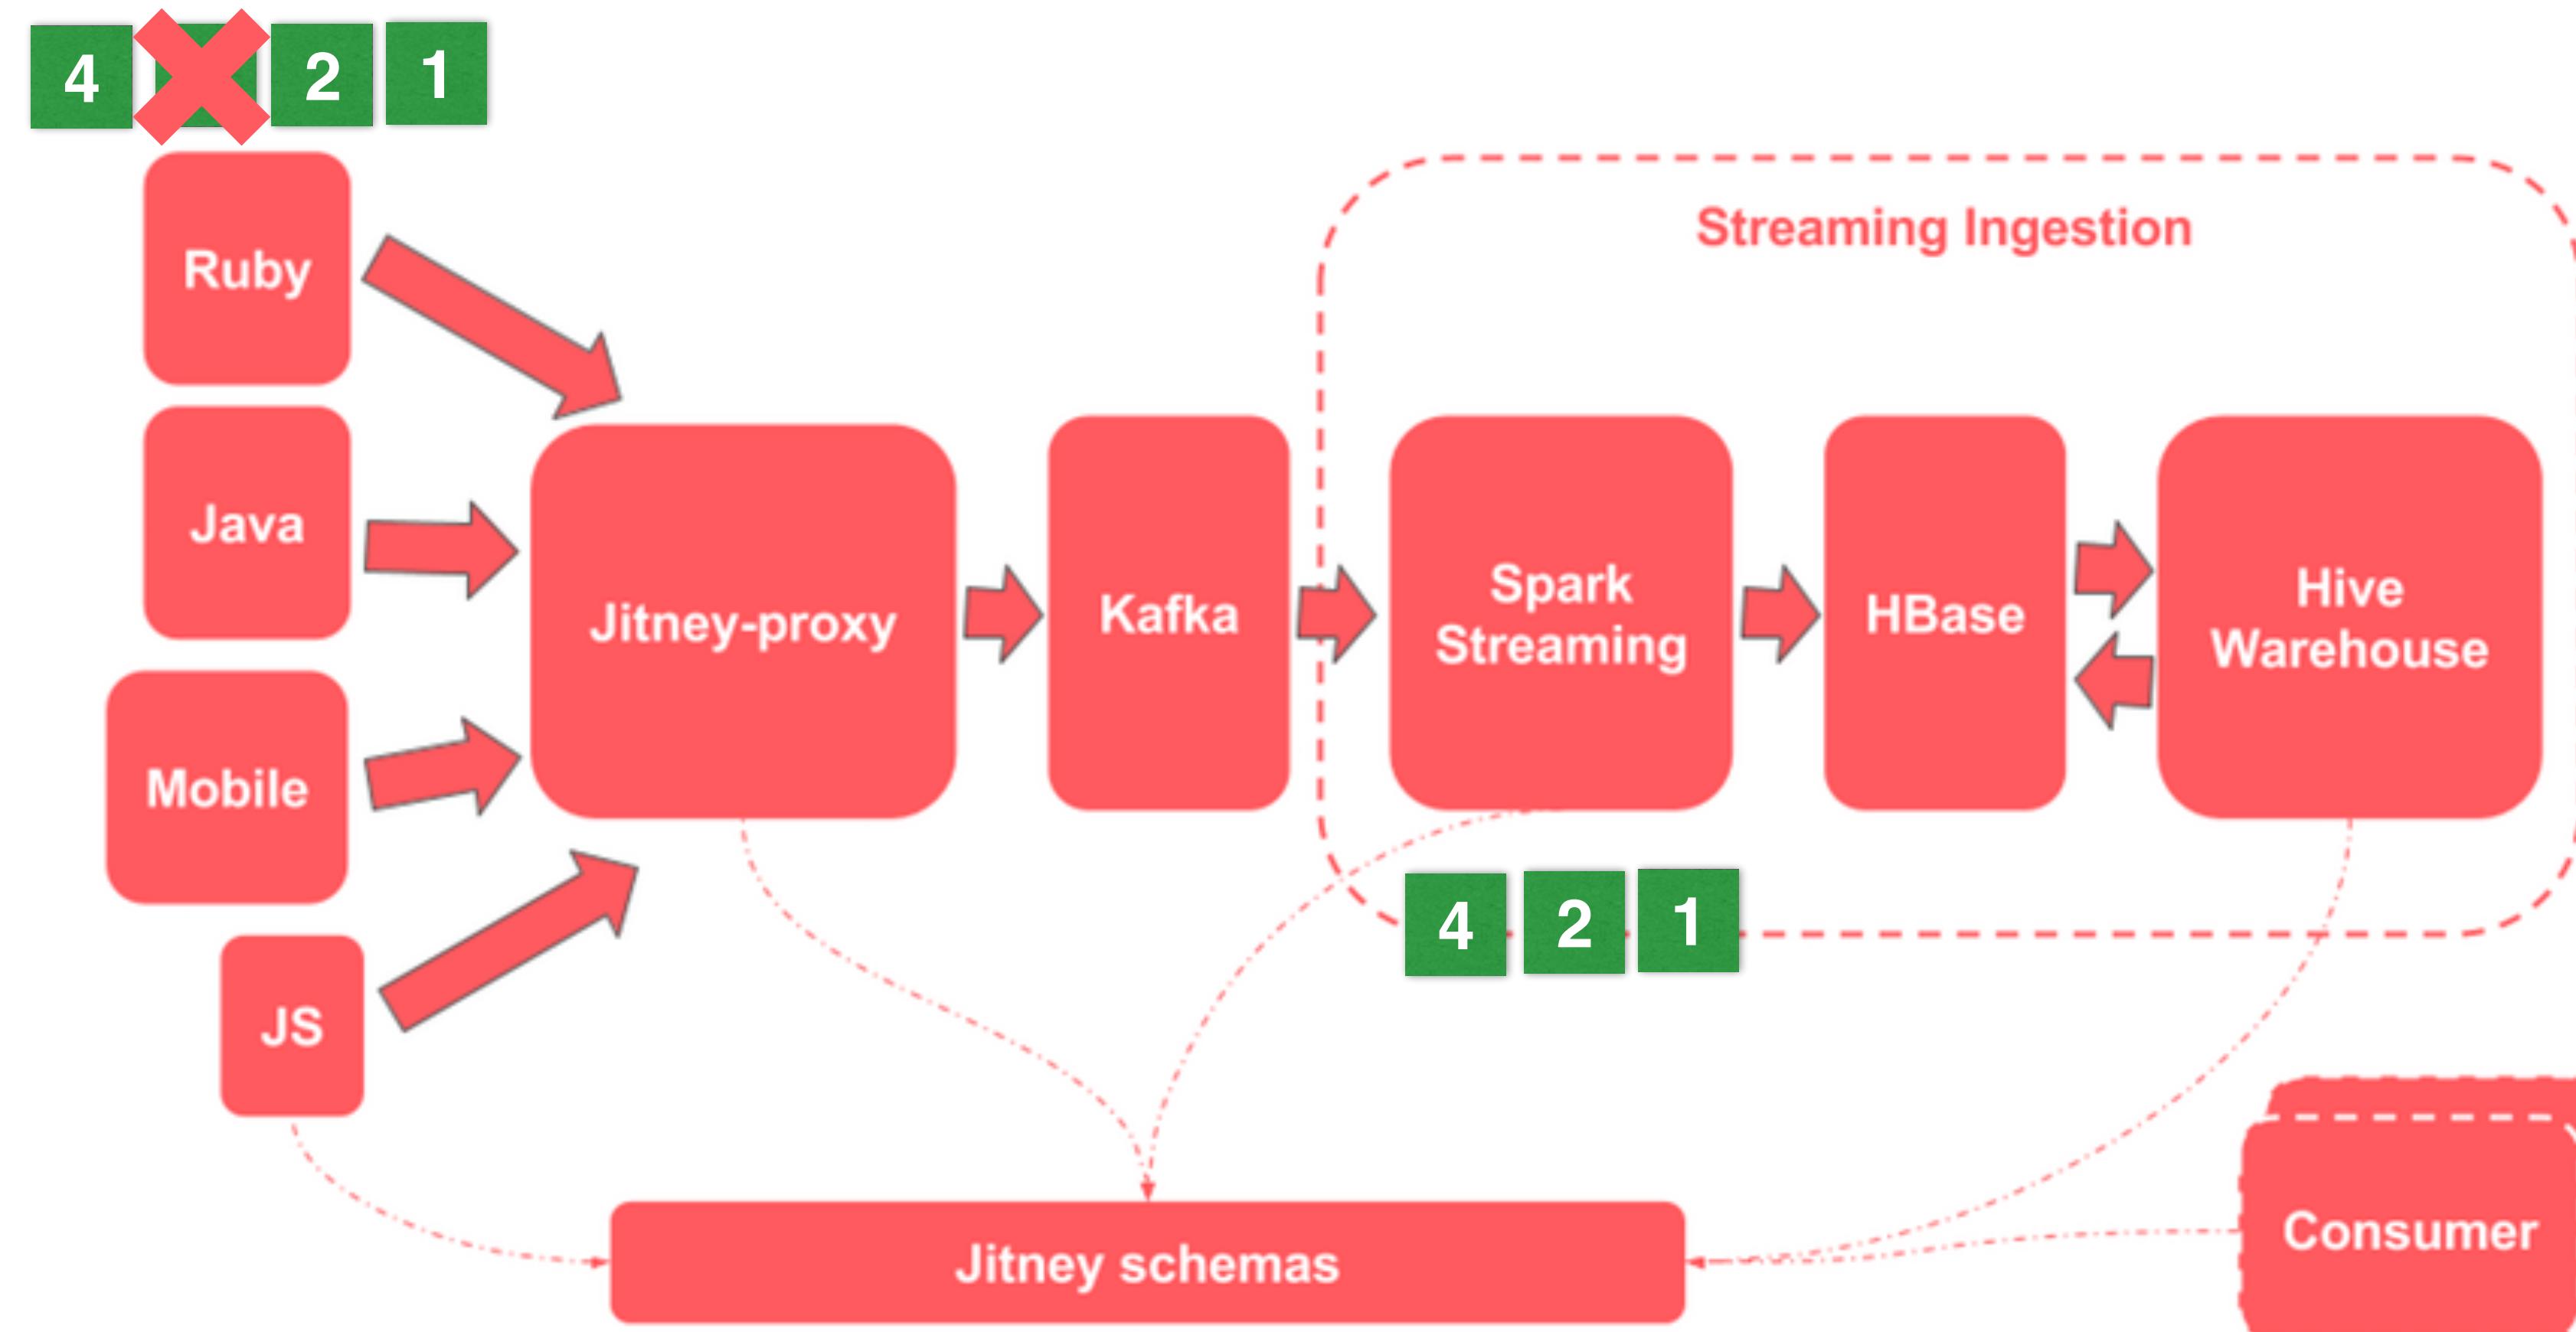
- Stack: Jitney, Spark Streaming, HBase, HDFS
- Spark Streaming 1.5 with Kafka “direct” connect
- Process 1 minute batches
- Write to HBase after deserializing with the right schema
- Dump data to HDFS every hour (with dedup) and add a Hive partition
- But live data can be queried via “current” partition

Data Ingestion Pipeline

end to end

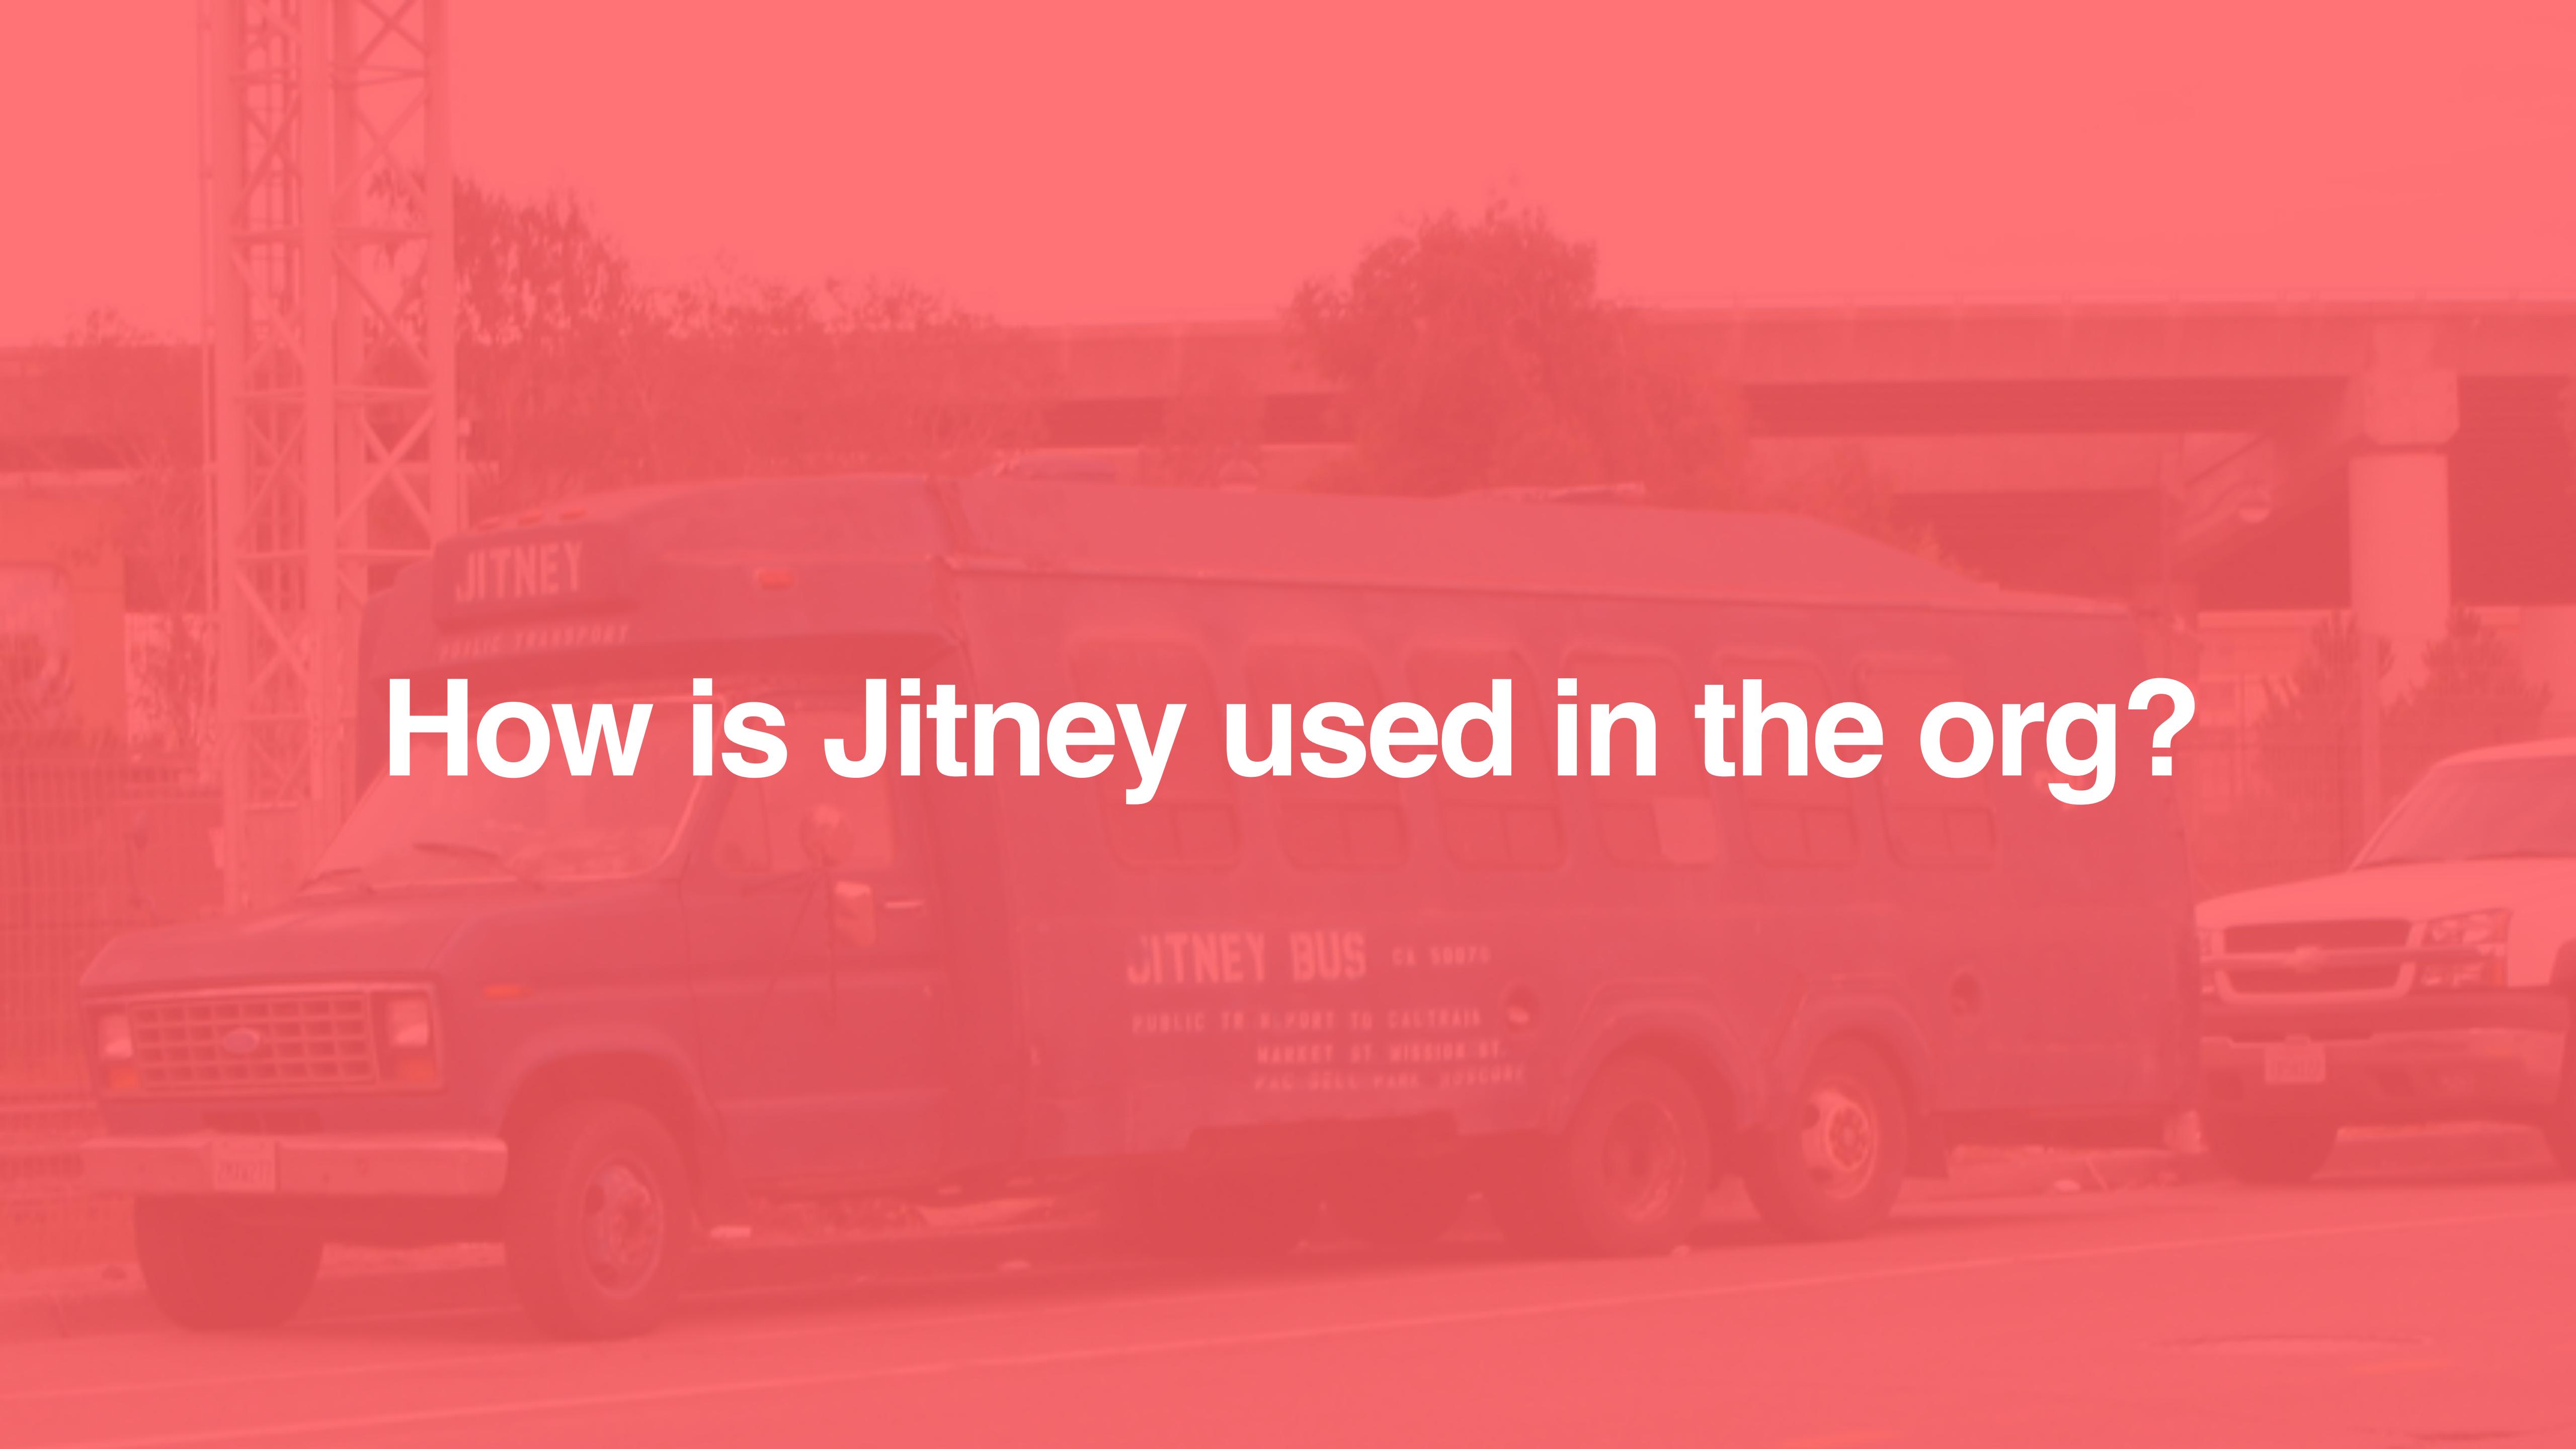


Audit

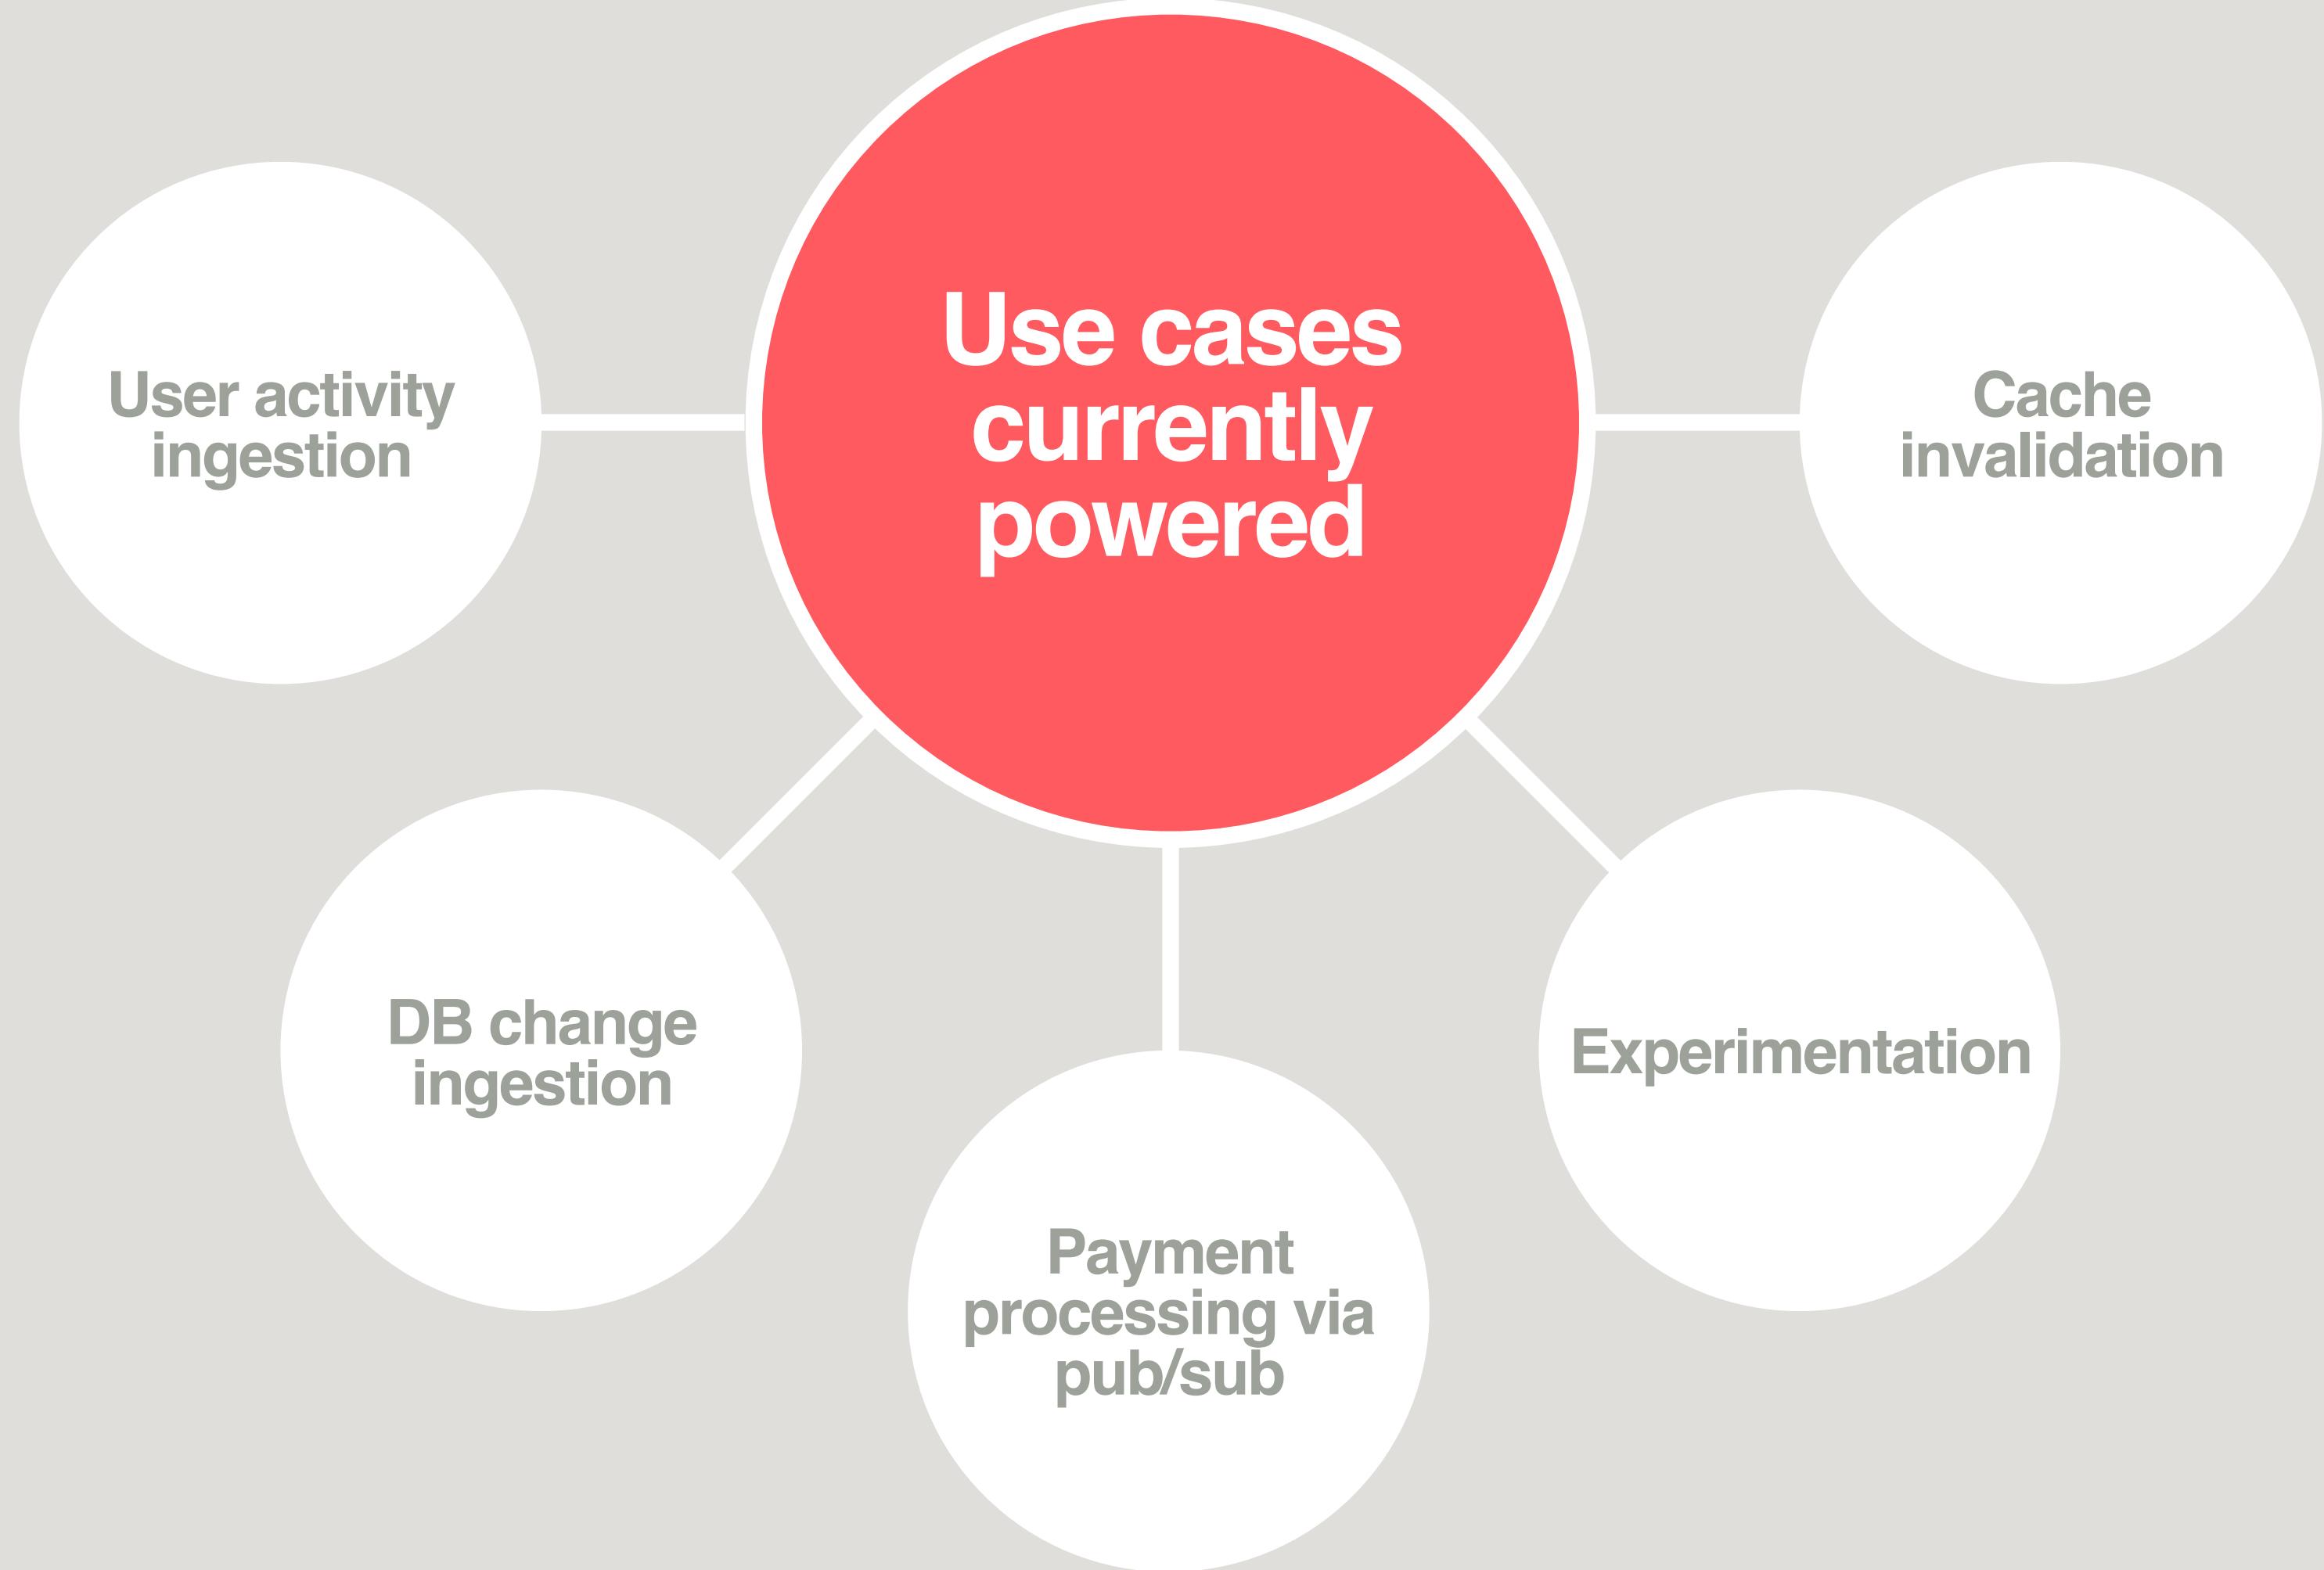


Event Schema for Audit Metadata

```
{ [ ]  
  "eventSchema": "com.airbnb.jitney.event.di_canary:CanaryEvent:1.0.0",  
  "id": "a941b1d3-67c8-41b5-af20-adb176bb53ee",  
  "tiers": [ [ ]  
    { [ ]  
      "id": "e2d43d8b-d23e-4bde-907f-7526ffd41a24",  
      "type": "jitney-ruby-client",  
      "hostname": "i-8c516c5f",  
      "ipAddress": "10.169.141.54",  
      "timestamp": 1456262927715,  
      "sequences": [ [ ]  
        { [ ]  
          "key": [ [ ]  
            "global"  
          ],  
          "value": 0,  
          "bytesSent": 141  
        }  
      ]  
    },  
    { [ ]  
    }  
  ]  
}
```

A vintage-style bus with "JITNEY" and "JITNEY BUS" signs. The bus is white with red and blue stripes. It has a sign on the front that reads "JITNEY BUS EX 5007 PUBLIC TO REPORT TO CALTRAIN MARKET AT MISSION ST. PAC-SIDE PARK SEPARATE".

How is Jitney used in the org?



Use cases currently powered

User activity
ingestion

Cache
invalidation

DB change
ingestion

Experimentation

Payment
processing via
pub/sub

Key take aways

1

Standardization!

2

Auditing Pipeline

3

Huge Advantage
for the organization

Join Airbnb



Learn More

<https://www.airbnb.com/careers/locations/beijing-china>